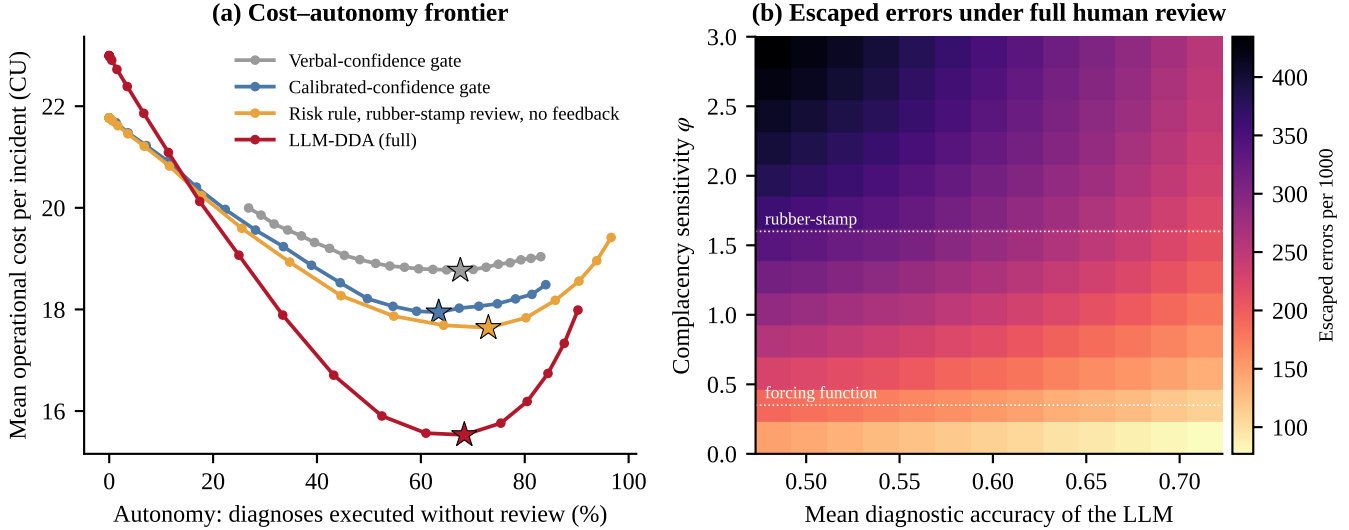
(a) Cost–autonomy frontier
Verbal-confidence gate
Calibrated-confidence gate
Risk rule, rubber-stamp review, no feedback
LLM-DDA (full)
Mean operational cost per incident (CU)
22
20
18
16
0
20
40
60
80
100
Autonomy: diagnoses executed without review (%)
(b) Escaped errors under full human review
Complacency sensitivity φ
3.0
2.5
2.0
1.5
1.0
0.5
0.0
rubber-stamp
forcing function
0.50
0.55
0.60
0.65
0.70
Mean diagnostic accuracy of the LLM
Escaped errors per 1000
400
350
300
250
200
150
100