Risk-conditioned delegation of LLM network diagnoses

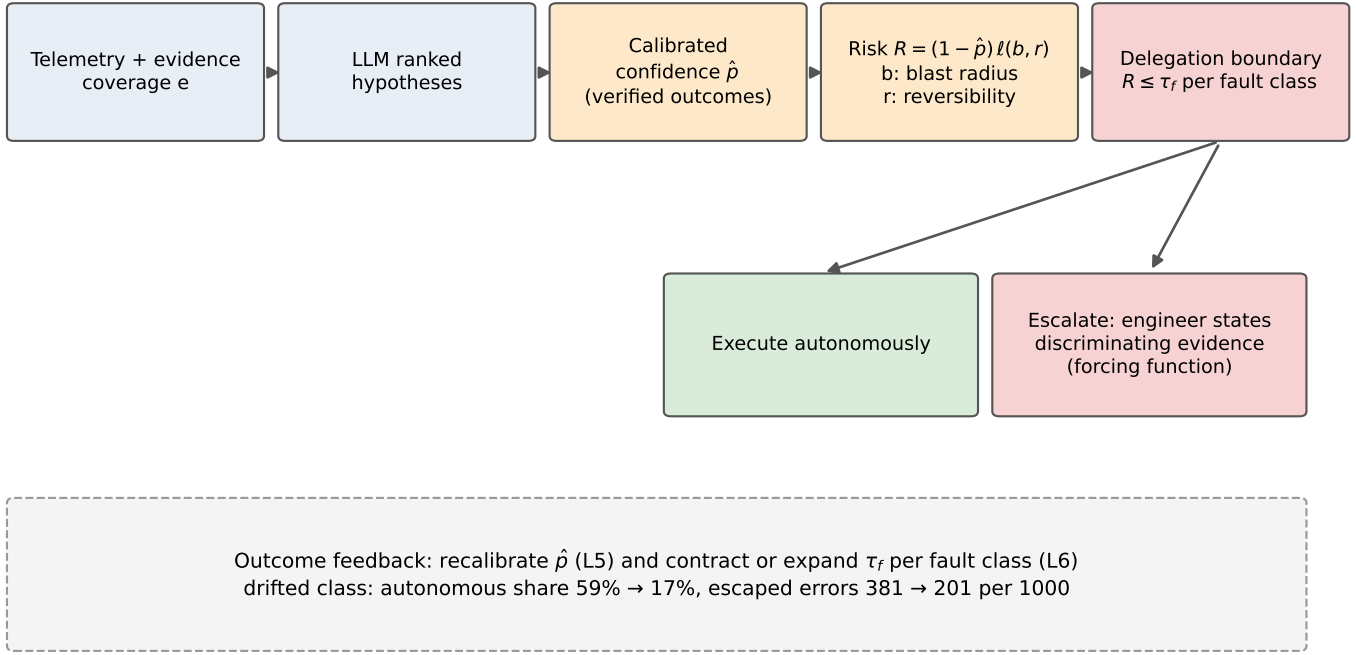

Simulation benchmark: 6,000-incident streams x 30 seeds, 6 fault classes,
overconfident LLM surrogate (self-confidence 0.95 vs accuracy 0.64), mid-stream drift.

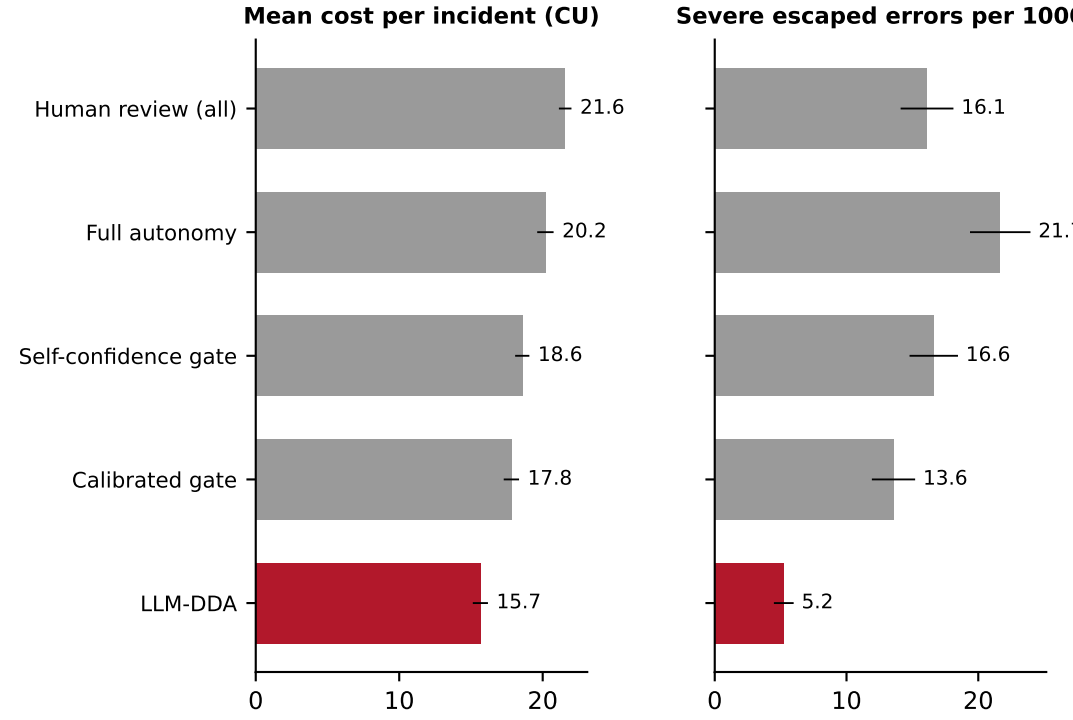

LLM-DDA vs calibrated gate: 12.1% lower cost, 61% fewer severe escapes
(30 seeds, paired Wilcoxon, Holm-adjusted p < 1e-5)