Risk-conditioned delegation of LLM network diagnoses

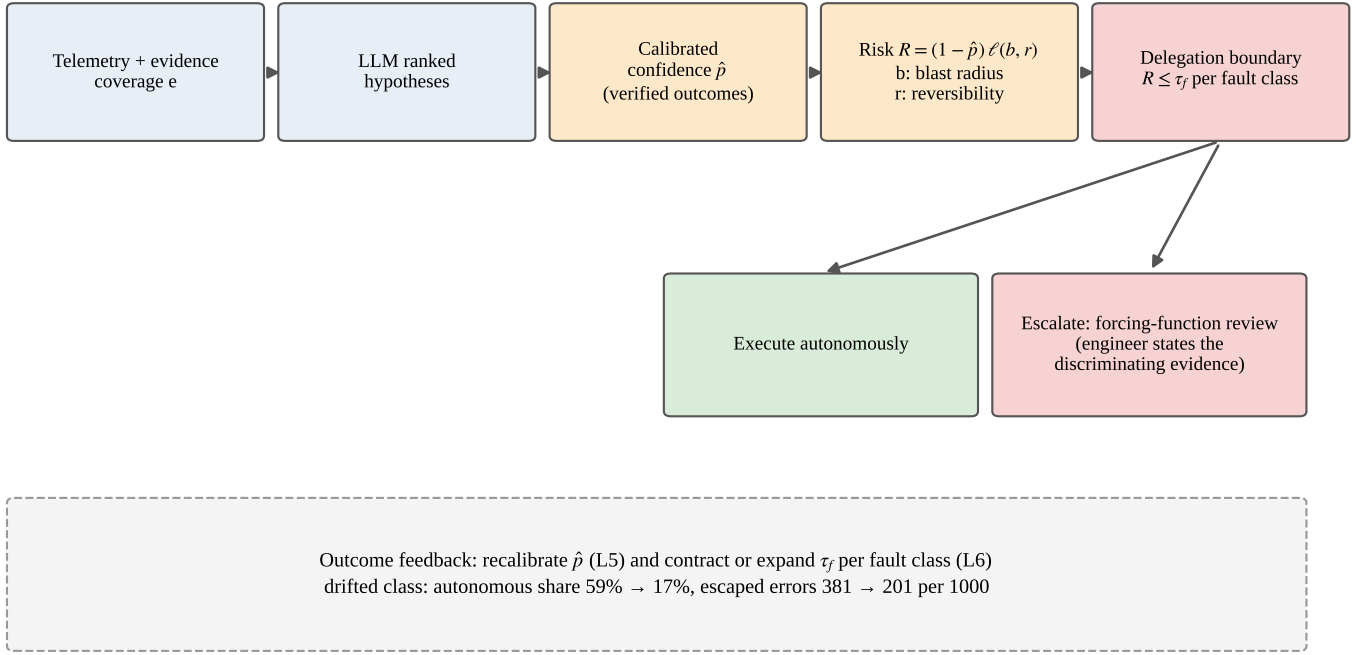

Simulation benchmark: 6,000-incident streams x 30 seeds, 6 fault classes,
overconfident LLM surrogate (verbal confidence 0.95 vs accuracy 0.64), mid-stream drift.

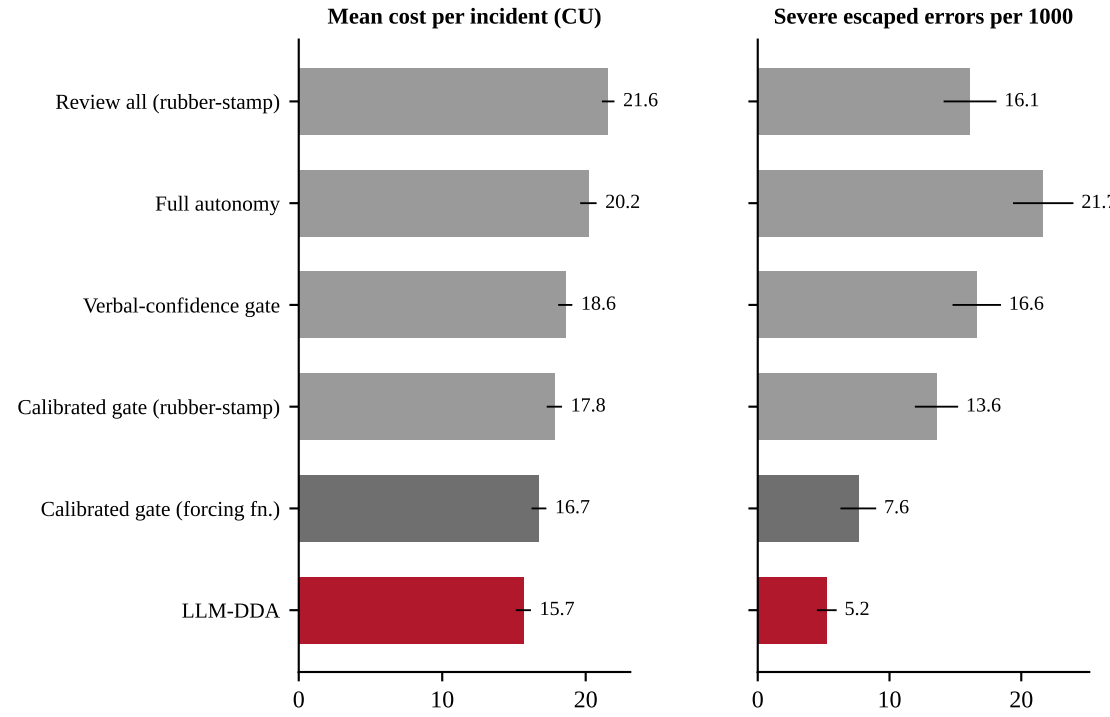

LLM-DDA vs calibrated gate: 12.1% lower cost and 61% fewer severe escapes (rubber-stamp review);
6.5% and 31% when the gate uses the same forcing-function review (30 seeds; all paired p < 1e-5)